

# Exploratory\_Data\_Analysis\_Assign

January 26, 2021

## 1 objective : Classify whether new patient can survive or not given x

Data-Set Link : <https://www.kaggle.com/gilsousa/habermans-survival-data-set>

Data-Set Description : The dataset contains cases from a study that was conducted between 1958 and 1970 at the University of Chicago's Billings Hospital on the survival of patients who had undergone surgery for breast cancer.

Number of Instances: 306

Number of Attributes: 4 (including the class attribute)

Attribute Information: f1 : Age of patient at time of operation (numerical) f2 : Patient's year of operation (year - 1900, numerical) f3 : Number of positive axillary nodes detected (numerical) f4 : Survival status (class attribute) 1 = the patient survived 5 years or longer 2 = the patient died within 5 year Missing Attribute Values: None

```
[1]: from google.colab import drive
drive.mount('/content/drive')
```

Mounted at /content/drive

```
[2]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import numpy as np
import warnings
warnings.filterwarnings("ignore")

df = pd.read_csv('/content/drive/MyDrive/AAIC/Assignments/4.Exploratory Data_
→Analysis on Haberman Dataset/haberman.csv')

col = ['Age of patient at time of operation',
       'Patients year of operation',
       'Number of positive axillary nodes detected',
       'survival_status_after_5_years']

df.columns = col
```

## 2 High level statistics of the dataset

```
[3]: # Number of Points , number of features, number of classes, data-points per
      ↪class.
print('Number of Points:',df.shape[0])
print('Number of features:',df.shape[1])
print('Number of classes:',set(df.iloc[:,-1]))
print('Number of data-points per class:\n',df['survival_status_after_5_years'].
      ↪value_counts())
```

```
Number of Points: 305
Number of features: 4
Number of classes: {1, 2}
Number of data-points per class:
1    224
2     81
Name: survival_status_after_5_years, dtype: int64
```

OBS: It seems like data is imbalanced

```
[4]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 305 entries, 0 to 304
Data columns (total 4 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   Age of patient at time of operation    305 non-null   int64
1   Patients year of operation             305 non-null   int64
2   Number of positive axillary nodes detected 305 non-null   int64
3   survival_status_after_5_years          305 non-null   int64
dtypes: int64(4)
memory usage: 9.7 KB
```

```
[5]: df.describe()
```

```
[5]:      Age of patient at time of operation  ...  survival_status_after_5_years
count      305.000000  ...      305.000000
mean       52.531148  ...      1.265574
std        10.744024  ...      0.442364
min        30.000000  ...      1.000000
25%        44.000000  ...      1.000000
50%        52.000000  ...      1.000000
75%        61.000000  ...      2.000000
max        83.000000  ...      2.000000
```


```
[8 rows x 4 columns]
```

### 3 Perform Univariate analysis(PDF, CDF, Boxplot, Violin plots) to understand which features are useful towards classification.


```
[6]: df.columns
```

```
[6]: Index(['Age of patient at time of operation', 'Patients year of operation',  
        'Number of positive axillary nodes detected',  
        'survival_status_after_5_years'],  
        dtype='object')
```

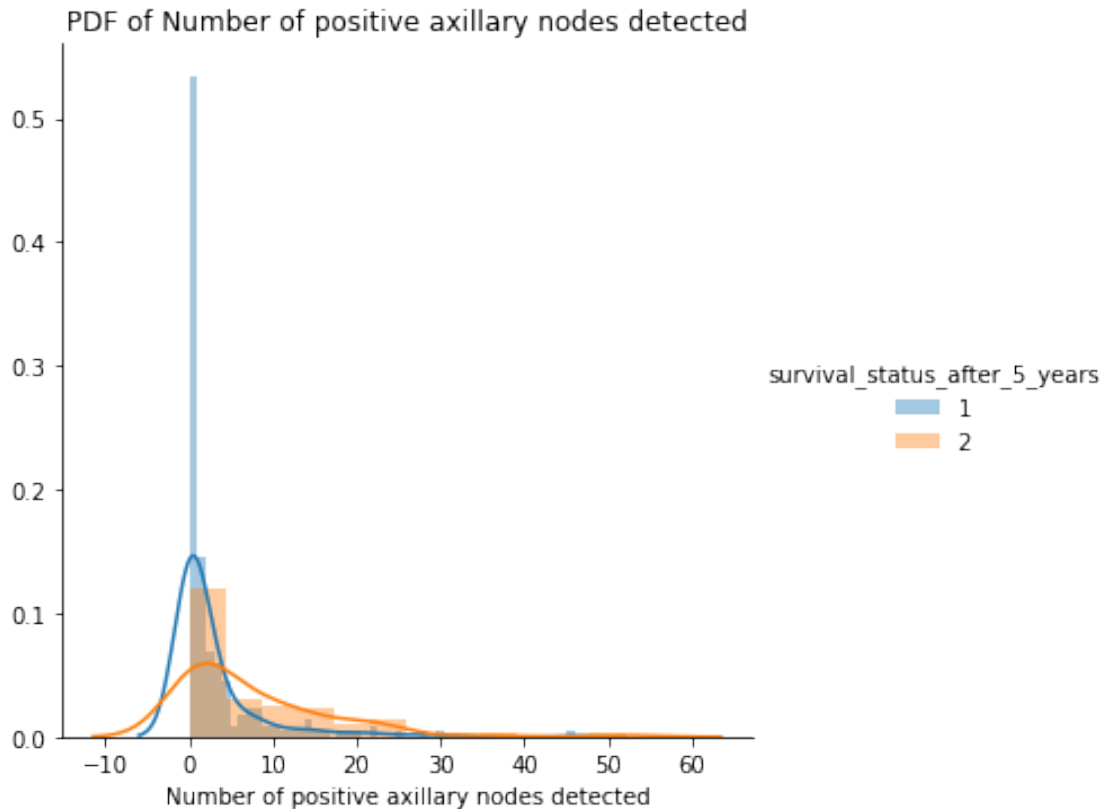
```
[7]: x = ['Age of patient at time of operation',  
        'Patients year of operation',  
        'Number of positive axillary nodes detected']  
  
for col in x:  
    sns.FacetGrid(df, hue="survival_status_after_5_years", height=5) \  
        .map(sns.distplot, col) \  
        .add_legend();  
    plt.title('PDF of {0}'.format(col))  
    plt.show();  
    print('*'*50)
```



\*\*\*\*\*



\*\*\*\*\*



\*\*\*\*\*

OBS: Here all the features are overlapping with the class labels, i.e the pdf plots of each feature w.r.t class labels are found to overlap their distribution, But for the feature " Number of positive axillary nodes detected ", after the count 5 there is a small gap b/w the distributions.

```
[15]: df.columns
```

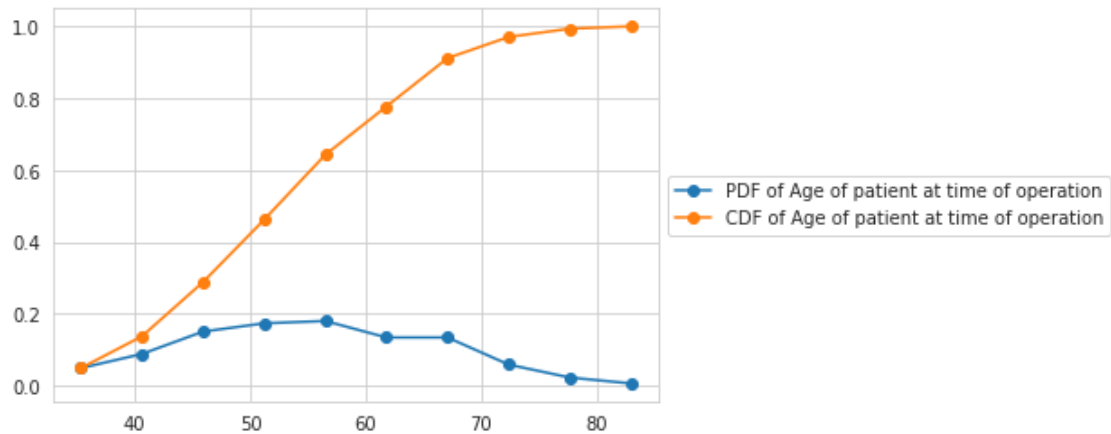
```
[15]: Index(['Age of patient at time of operation', 'Patients year of operation',
         'Number of positive axillary nodes detected',
         'survival_status_after_5_years'],
        dtype='object')
```

```
[55]: counts, bin_edges = np.histogram(df['Age of patient at time of operation'],
        ↪bins=10,
        density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,'-o',label='PDF of Age of patient at time of_
        ↪operation')
```

```
plt.plot(bin_edges[1:], cdf, '-o', label='CDF of Age of patient at time of_
→operation')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.show()
```


```
[0.04918033 0.08852459 0.15081967 0.17377049 0.18032787 0.13442623
 0.13442623 0.05901639 0.02295082 0.00655738]
[30.  35.3 40.6 45.9 51.2 56.5 61.8 67.1 72.4 77.7 83. ]
```



```
[56]: counts, bin_edges = np.histogram(df['Patients year of operation'], bins=10,
                                     density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf, '-o', label='PDF of Patients year of operation')
plt.plot(bin_edges[1:], cdf, '-o', label='CDF of Patients year of operation')
plt.legend(loc='center left', bbox_to_anchor=(1, 0.5))
plt.show()
```

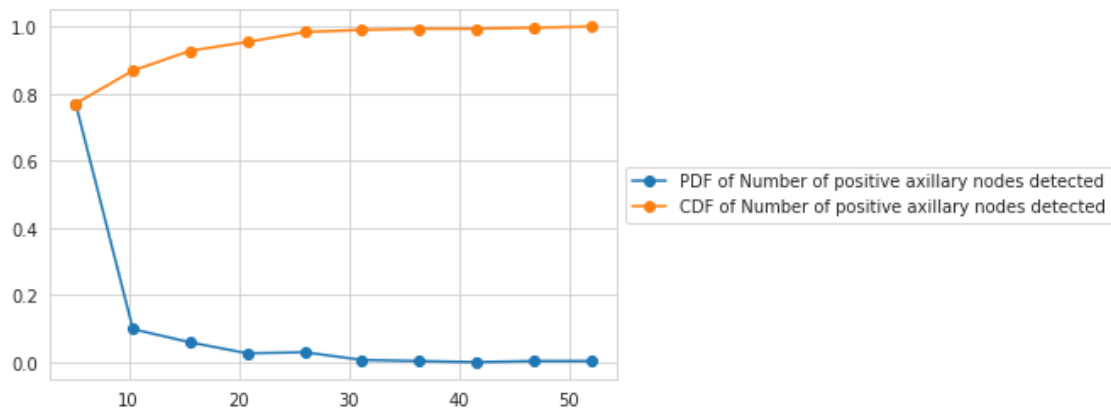
```
[0.20655738 0.09180328 0.0852459  0.07540984 0.09836066 0.09836066
 0.09180328 0.09180328 0.08196721 0.07868852]
[58.  59.1 60.2 61.3 62.4 63.5 64.6 65.7 66.8 67.9 69. ]
```



```
[57]: counts, bin_edges = np.histogram(df['Number of positive axillary nodes_
    →detected'], bins=10,
                                     density = True)

pdf = counts/(sum(counts))
print(pdf);
print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,'-o',label='PDF of Number of positive axillary nodes_
    →detected')
plt.plot(bin_edges[1:], cdf,'-o',label='CDF of Number of positive axillary_
    →nodes detected')
plt.legend(loc='center left',bbox_to_anchor=(1, 0.5))
plt.show()
```

```
[0.7704918  0.09836066 0.05901639 0.02622951 0.0295082  0.00655738
 0.00327869 0.          0.00327869 0.00327869]
[ 0.   5.2 10.4 15.6 20.8 26.   31.2 36.4 41.6 46.8 52. ]
```



```

[58]: counts, bin_edges = np.histogram(df['Age of patient at time of operation'],
    ↪ bins=10,
    density = True)

pdf = counts/(sum(counts))
#print(pdf);
#print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,'-o',label='PDF of Age of patient at time of
    ↪ operation')
plt.plot(bin_edges[1:], cdf,'-o',label='CDF of Age of patient at time of
    ↪ operation')

counts, bin_edges = np.histogram(df['Patients year of operation'], bins=10,
    density = True)

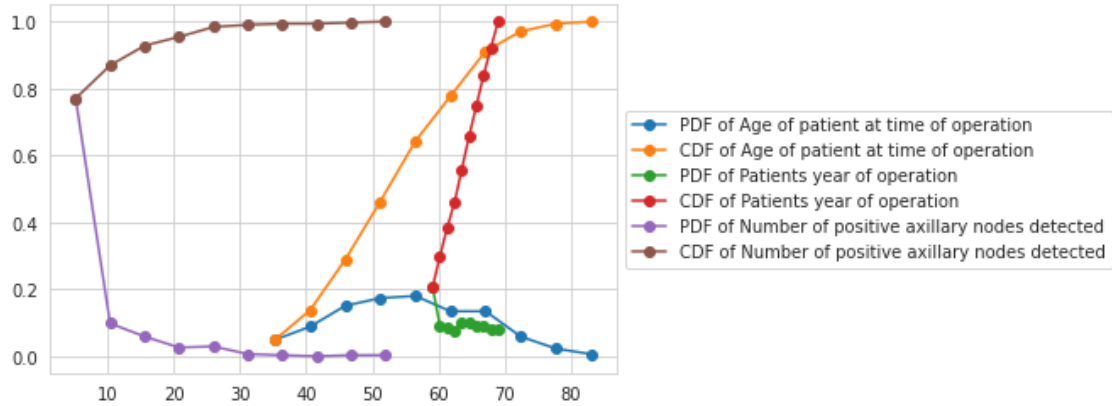
pdf = counts/(sum(counts))
#print(pdf);
#print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,'-o',label='PDF of Patients year of operation')
plt.plot(bin_edges[1:], cdf,'-o',label='CDF of Patients year of operation')

counts, bin_edges = np.histogram(df['Number of positive axillary nodes
    ↪ detected'], bins=10,
    density = True)

pdf = counts/(sum(counts))
#print(pdf);
#print(bin_edges)
cdf = np.cumsum(pdf)
plt.plot(bin_edges[1:],pdf,'-o',label='PDF of Number of positive axillary nodes
    ↪ detected')
plt.plot(bin_edges[1:], cdf,'-o',label='CDF of Number of positive axillary
    ↪ nodes detected')
plt.legend(loc='center left',bbox_to_anchor=(1, 0.5))
plt.show();

```





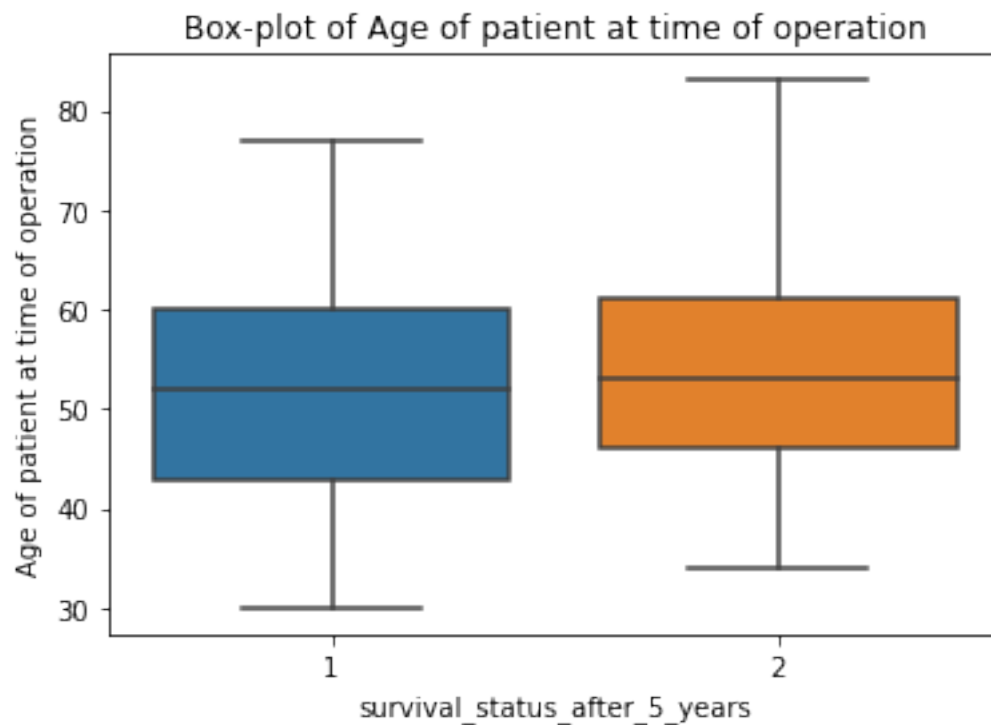
```
[23]: df.columns
```

```
[23]: Index(['Age of patient at time of operation', 'Patients year of operation',
          'Number of positive axillary nodes detected',
          'survival_status_after_5_years'],
          dtype='object')
```


```
[41]: from prettytable import PrettyTable
t = PrettyTable(['Features', 'Medians', 'Quantiles' , '90th Percentiles' ,
→ 'Median Absolute Deviation'])
#print(clf.best_estimator_,'\n',auc_test)
t.add_row(['Age of patient at time of operation',\
          (np.median(df['Age of patient at time of operation'])),\
          (np.percentile(df['Age of patient at time of operation'],np.
→ arange(0, 100, 25))),\
          (np.percentile(df['Age of patient at time of operation'],90)),\
          (robust.mad(df['Age of patient at time of operation']))]
t.add_row(['Patients year of operation',\
          (np.median(df['Patients year of operation'])),\
          (np.percentile(df['Patients year of operation'],np.arange(0, 100,
→ 25))),\
          (np.percentile(df['Patients year of operation'],90)),\
          (robust.mad(df['Patients year of operation']))]
t.add_row(['Number of positive axillary nodes detected',\
          (np.median(df['Number of positive axillary nodes detected'])),\
          (np.percentile(df['Number of positive axillary nodes detected'], np.
→ arange(0, 100, 25))),\
          (np.percentile(df['Number of positive axillary nodes detected'],
→ 90)),\
          (robust.mad(df['Number of positive axillary nodes detected']))]
print(t)
```

Features		Medians	Quantiles
90th Percentiles	Median Absolute Deviation		
Age of patient at time of operation		52.0	[30. 44. 52. 61.]
67.0	11.860817748044816		
Patients year of operation		63.0	[58. 60. 63. 66.]
67.0	4.447806655516806		
Number of positive axillary nodes detected		1.0	[0. 0. 1. 4.]
13.0	1.482602218505602		

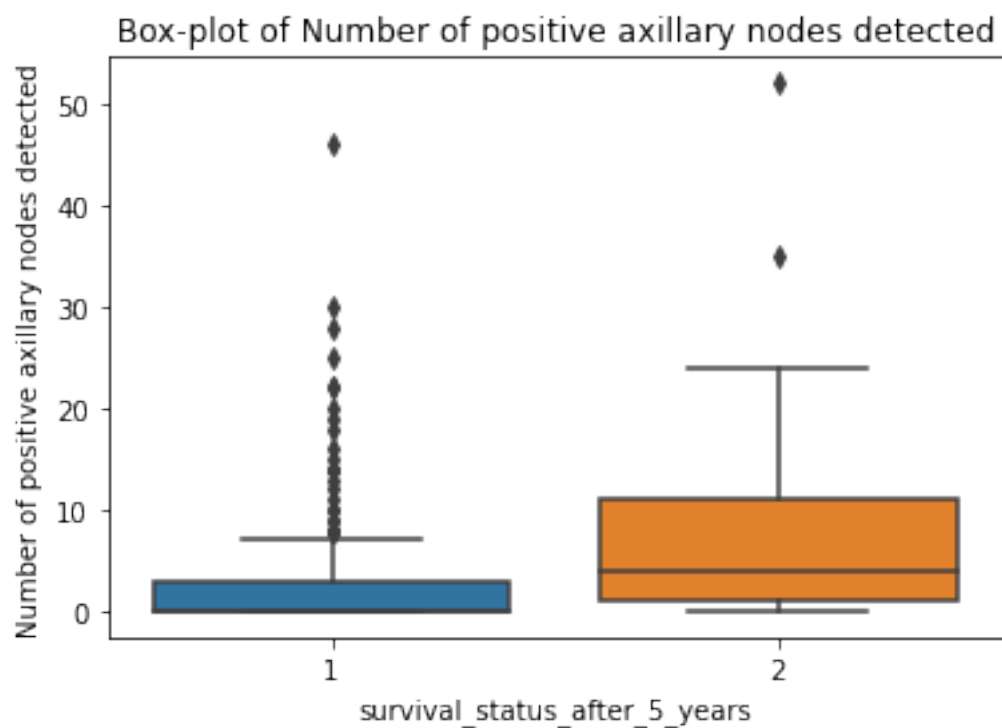
```
[1]:
[8]: for col in x:
      sns.boxplot(x='survival_status_after_5_years',y = col, data=df)
      plt.title('Box-plot of {0}'.format(col))
      plt.show()
      print('*'*50)
```



\*\*\*\*\*



\*\*\*\*\*




\*\*\*\*\*

OBS: we can observe some difference in the box-plot curve for the feature " Number of positive axillary nodes detected ". I think this feature can be used in classification models


```
[ ]:
[9]: for col in x:
      sns.violinplot(x='survival_status_after_5_years',y = col, data=df)
      plt.title('violin-plot of {0}'.format(col))
      plt.show()
      print('*'*50)
```



\*\*\*\*\*



\*\*\*\*\*



\*\*\*\*\*

[ ]:

#### 4 Perform Bi-variate analysis (scatter plots, pair-plots) to see if combinations of features are useful in classification.


[10]: `df.columns`

[10]: `Index(['Age of patient at time of operation', 'Patients year of operation',  
'Number of positive axillary nodes detected',  
'survival_status_after_5_years'],  
dtype='object')`

[11]: `sns.set_style("whitegrid");  
sns.FacetGrid(df, hue="survival_status_after_5_years", height=8) \  
 .map(plt.scatter, 'Age of patient at time of operation', 'Patients year of_  
→operation') \  
 .add_legend();  
  
plt.show();`




```
[12]: sns.set_style("whitegrid");
sns.FacetGrid(df, hue="survival_status_after_5_years", height=8) \
    .map(plt.scatter, 'Age of patient at time of operation', 'Number of positive_
    →axillary nodes detected') \
    .add_legend();
plt.show();
```

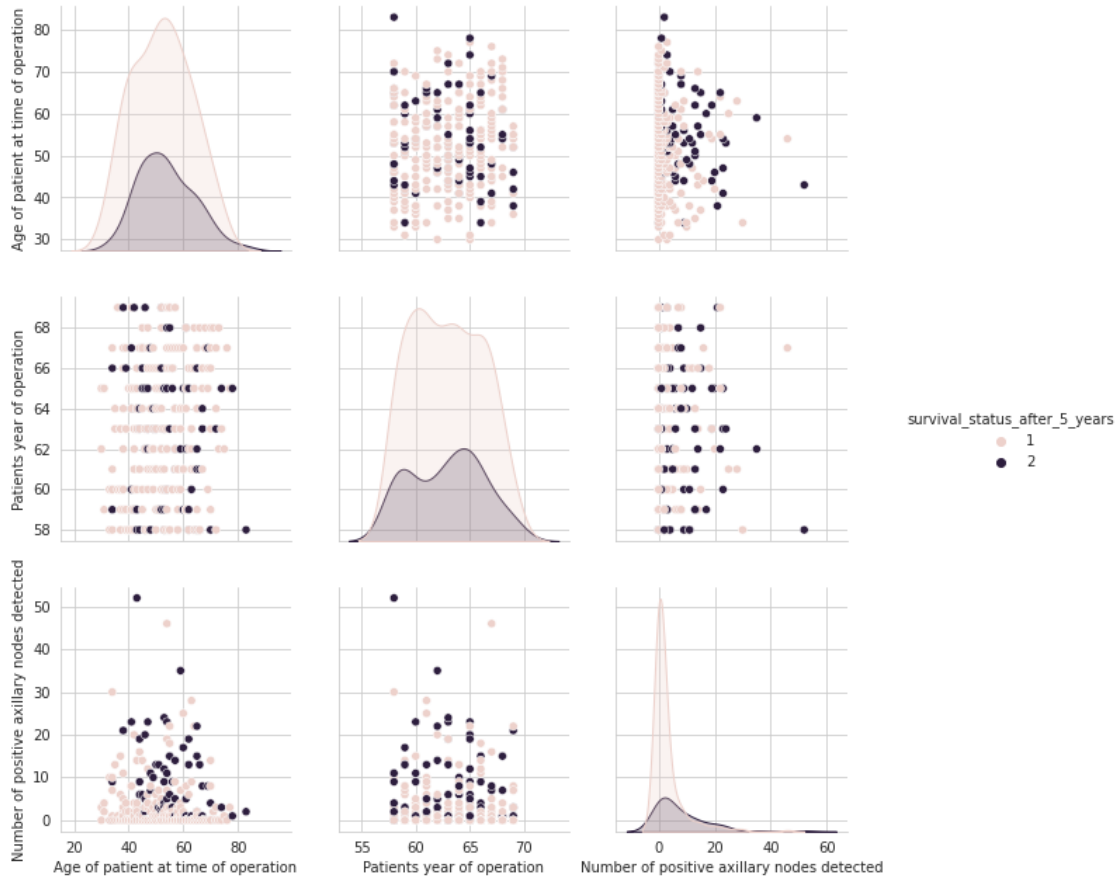


```
[13]: sns.set_style("whitegrid");
sns.FacetGrid(df, hue="survival_status_after_5_years", height=8) \
    .map(plt.scatter, 'Patients year of operation', 'Number of positive axillary_
    ↪nodes detected') \
    .add_legend();
plt.show();
```





```
[14]: sns.set_style("whitegrid");  
sns.pairplot(df, hue="survival_status_after_5_years", height=3);  
plt.show()
```



OBS: There is an overlap of data points in this scatter-plots. We cannot find a line/curve which divides this overlap until unless projected into 3-d or n-d.

[ ]:

[ ]:

[ ]:

[ ]: