

9_Implement_SGD_Classifier_with_Logloss_and_L2_regularization_U

November 2, 2020

1 Implement SGD Classifier with Logloss and L2 regularization Using SGD without using sklearn

There will be some functions that start with the word "grader" ex: grader_weights(), grader_sigmoid(), grader_logloss() etc, you should not change those function definition. Every Grader function has to return True.

Importing packages

```
[1]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler
from sklearn import linear_model
```

Creating custom dataset

```
[2]: # please don't change random_state
X, y = make_classification(n_samples=50000, n_features=15, n_informative=10,
    →n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7,
    →random_state=15)
# make_classification is used to create custom dataset
# Please check this link (https://scikit-learn.org/stable/modules/generated/sklearn.datasets.make\_classification.html) for more details
```

```
[3]: X.shape, y.shape
```

```
[3]: ((50000, 15), (50000,))
```

Splitting data into train and test

```
[4]: #please don't change random state
x_train, x_test, y_train, y_test = train_test_split(X, y, test_size=0.25,
    →random_state=15)
```

```
[5]: # Standardizing the data.
scaler = StandardScaler()
x_train = scaler.fit_transform(x_train)
```

```
x_test = scaler.transform(x_test)
```

```
[6]: x_train.shape, y_train.shape, x_test.shape, y_test.shape
```

```
[6]: ((37500, 15), (37500,), (12500, 15), (12500,))
```

```
[6]:
```

2 SGD classifier

```
[7]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the constant, invscaling or adaptive schedules.

clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log',
    ↪random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
# Please check this documentation (https://scikit-learn.org/stable/modules/generated/sklearn.linear\_model.SGDClassifier.html)
```

```
[7]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
[8]: clf.fit(X=x_train, y=y_train) # fitting our model
```

```
-- Epoch 1
```

```
Norm: 0.70, NNZs: 15, Bias: -0.501317, T: 37500, Avg. loss: 0.552526
```

```
Total training time: 0.01 seconds.
```

```
-- Epoch 2
```

```
Norm: 1.04, NNZs: 15, Bias: -0.752393, T: 75000, Avg. loss: 0.448021
```

```
Total training time: 0.03 seconds.
```

```
-- Epoch 3
```

```
Norm: 1.26, NNZs: 15, Bias: -0.902742, T: 112500, Avg. loss: 0.415724
```

```
Total training time: 0.04 seconds.
```

```
-- Epoch 4
```

```
Norm: 1.43, NNZs: 15, Bias: -1.003816, T: 150000, Avg. loss: 0.400895
```

```
Total training time: 0.05 seconds.
```

```
-- Epoch 5
```

```
Norm: 1.55, NNZs: 15, Bias: -1.076296, T: 187500, Avg. loss: 0.392879
```

```
Total training time: 0.06 seconds.
```

```
-- Epoch 6
```

```
Norm: 1.65, NNZs: 15, Bias: -1.131077, T: 225000, Avg. loss: 0.388094
```

```

Total training time: 0.07 seconds.
-- Epoch 7
Norm: 1.73, NNZs: 15, Bias: -1.171791, T: 262500, Avg. loss: 0.385077
Total training time: 0.08 seconds.
-- Epoch 8
Norm: 1.80, NNZs: 15, Bias: -1.203840, T: 300000, Avg. loss: 0.383074
Total training time: 0.09 seconds.
-- Epoch 9
Norm: 1.86, NNZs: 15, Bias: -1.229563, T: 337500, Avg. loss: 0.381703
Total training time: 0.10 seconds.
-- Epoch 10
Norm: 1.90, NNZs: 15, Bias: -1.251245, T: 375000, Avg. loss: 0.380763
Total training time: 0.11 seconds.
-- Epoch 11
Norm: 1.94, NNZs: 15, Bias: -1.269044, T: 412500, Avg. loss: 0.380084
Total training time: 0.12 seconds.
-- Epoch 12
Norm: 1.98, NNZs: 15, Bias: -1.282485, T: 450000, Avg. loss: 0.379607
Total training time: 0.13 seconds.
-- Epoch 13
Norm: 2.01, NNZs: 15, Bias: -1.294386, T: 487500, Avg. loss: 0.379251
Total training time: 0.14 seconds.
-- Epoch 14
Norm: 2.03, NNZs: 15, Bias: -1.305805, T: 525000, Avg. loss: 0.378992
Total training time: 0.15 seconds.
Convergence after 14 epochs took 0.15 seconds

```

```

[8]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)

```

```

[9]: clf.coef_, clf.coef_.shape, clf.intercept_
    #clf.coef_ will return the weights
    #clf.coef_.shape will return the shape of weights
    #clf.intercept_ will return the intercept term

```

```

[9]: (array([[ -0.89007184,  0.63162363, -0.07594145,  0.63107107, -0.38434375,
    0.93235243, -0.89573521, -0.07340522,  0.40591417,  0.4199991 ,
    0.24722143,  0.05046199, -0.08877987,  0.54081652,  0.06643888]]),
    (1, 15),
    array([-1.30580538]))

```

```

# This is formatted as code

```

2.1 Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

- Initialize the weight_vector and intercept term to zeros (Write your code in def initialize_weights())
- Create a loss function (Write your code in def logloss())

$logloss = -1 * \frac{1}{n} \sum_{foreach Y_t, Y_{pred}} (Y_t \log_{10}(Y_{pred}) + (1 - Y_t) \log_{10}(1 - Y_{pred}))$ - for each epoch:

- for each batch of data points in train: (keep batch size=1)

- calculate the gradient of loss function w.r.t each weight in weight vector (write your code in def grad_weight())

$$\frac{dw}{dt} = x_n(y_n - ((w^{(t)})^T x_n + b^{(t)})) - \frac{1}{N} w^{(t)}$$

- Calculate the gradient of the intercept (write your code in def grad_bias())

$$\frac{db}{dt} = y_n - ((w^{(t)})^T x_n + b^{(t)})$$

- Update weights and intercept (check the equation number 32 in the above mentioned [link](#))
 $w^{(t+1)} = w^{(t)} + (dw^{(t)})$

$$b^{(t+1)} = b^{(t)} + (db^{(t)})$$

- calculate the log loss for train and test with the updated weights (you can check the python code in the [link](#))
- And if you wish, you can compare the previous loss and the current loss, if it is not updating you can stop the training
- append this loss in the list (this will be used to see how loss is changing for each epoch)

Initialize weights

```
[10]: def initialize_weights(dim):
    ''' In this function, we will initialize our weights and bias'''
    #initialize the weights to zeros array of (1,dim) dimensions
    #you use zeros_like function to initialize zero, check this link https://
    →docs.scipy.org/doc/numpy/reference/generated/numpy.zeros_like.html
    #initialize bias to zero
    b = 0
    w = np.zeros_like(dim)

    return w,b
```

```
[10]:
[11]: dim=x_train[0]
w,b = initialize_weights(dim)
print('w =',(w))
print('b =',str(b))
```

```
w = [0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0. 0.]
b = 0
```

Grader function - 1

```
[12]: dim=x_train[0]
w,b = initialize_weights(dim)
def grader_weights(w,b):
    assert((len(w)==len(dim)) and b==0 and np.sum(w)==0.0)
    return True
grader_weights(w,b)
```

[12]: True

Compute sigmoid

$$\text{sigmoid}(z) = 1/(1 + \exp(-z))$$

```
[13]: def sigmoid(z):
    ''' In this function, we will return sigmoid of z'''
    # compute sigmoid(z) and return

    return 1/(1+np.exp(-z))
```

Grader function - 2

```
[14]: def grader_sigmoid(z):
    val=sigmoid(z)
    assert(val==0.8807970779778823)
    return True
grader_sigmoid(2)
```

[14]: True

Compute loss

$$\text{logloss} = -1 * \frac{1}{n} \sum_{\text{foreach } Y_t, Y_{\text{pred}}} (Y_t \log_{10}(Y_{\text{pred}}) + (1 - Y_t) \log_{10}(1 - Y_{\text{pred}}))$$

```
[15]: import math
def logloss(y_true,y_pred):
    '''In this function, we will compute log loss '''
    l=0
    #print(y_pred)
    for i in range(len(y_true)):
        l+= ( (y_true[i]*(math.log10(y_pred[i]))) + ((1-y_true[i])*(math.
→log10(1-y_pred[i]))) )

    loss = (-1/len(y_true))*l
    return loss
```

Grader function - 3

```
[16]: def grader_logloss(true,pred):
    loss=logloss(true,pred)
    assert(loss==0.07644900402910389)
    return True
true=[1,1,0,1,0]
pred=[0.9,0.8,0.1,0.8,0.2]
grader_logloss(true,pred)
```

[16]: True

Compute gradient w.r.to 'w'
 $dw^{(t)} = x_n(y_n((w^{(t)})^T x_n + b^t)) - \frac{1}{N}w^{(t)}$

[16]:

```
[17]: def gradient_dw(x,y,w,b,alpha,N):  
    '''In this function, we will compute the gradient w.r.to w '''  
    #print(x.shape,w.shape)  
    z = np.dot(x,w) + b  
    #print(z)  
    #print((y-sigmoid(z)),x)  
    a = ((y-sigmoid(z))*x)  
    #print(a)  
    b = ((alpha/N)*w)  
    dw = a-b  
    return dw
```

Grader function - 4

```
[18]: def grader_dw(x,y,w,b,alpha,N):  
    grad_dw=gradient_dw(x,y,w,b,alpha,N)  
    assert(np.sum(grad_dw)==2.613689585)  
    return True  
grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.  
→14783286,  
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,  
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])  
grad_y=0  
grad_w,grad_b=initialize_weights(grad_x)  
alpha=0.0001  
N=len(x_train)  
grader_dw(grad_x,grad_y,grad_w,grad_b,alpha,N)
```

[18]: True

Compute gradient w.r.to 'b'
 $\delta b^{(t)} = y_n - ((w^{(t)})^T x_n + b^{(t)})$

```
[19]: def gradient_db(x,y,w,b):  
    '''In this function, we will compute gradient w.r.to b '''  
    z = ( np.dot(x,w) + b )  
    db = (y-(sigmoid(z)))  
    return (db)
```

Grader function - 5

```
[20]: def grader_db(x,y,w,b):  
    grad_db=gradient_db(x,y,w,b)  
    assert(grad_db==-0.5)  
    return True
```

```

grad_x=np.array([-2.07864835,  3.31604252, -0.79104357, -3.87045546, -1.
→14783286,
                -2.81434437, -0.86771071, -0.04073287,  0.84827878,  1.99451725,
                3.67152472,  0.01451875,  2.01062888,  0.07373904, -5.54586092])
grad_y=0
grad_w,grad_b=initialize_weights(grad_x)
alpha=0.0001
N=len(x_train)
grader_db(grad_x,grad_y,grad_w,grad_b)

```

[20]: True

Implementing logistic regression

```

[21]: def pred(w,b, X):
        N = len(X)
        predict = []
        w = list(w)
        for i in range(N):
            z = np.dot(X[i],w) + b
            predict.append(sigmoid(z))
        return np.array(predict)
print(1-np.sum(y_train - pred(w,b,x_train))/len(x_train))
print(1-np.sum(y_test - pred(w,b,x_test))/len(x_test))

```

1.1978933333333333

1.19864

```

[22]: from tqdm import tqdm
from collections import OrderedDict
def train(x_train,y_train,x_test,y_test,epochs,alpha,eta0):
    ''' In this function, we will implement logistic regression'''
    #Here eta0 is learning rate
    #implement the code as follows
    # initialize the weights (call the initialize_weights(X_train[0]) function)
    # for every epoch
        # for every data point(X_train,y_train)
            #compute gradient w.r.to w (call the gradient_dw() function)
            #compute gradient w.r.to b (call the gradient_db() function)
            #update w, b
        # predict the output of x_train[for all data points in X_train] using
→w,b
        #compute the loss between predicted and actual values (call the loss
→function)
        # store all the train loss values in a list
        # predict the output of x_test[for all data points in X_test] using w,b
        #compute the loss between predicted and actual values (call the loss
→function)
        # store all the test loss values in a list

```

```

        # you can also compare previous loss and current loss, if loss is not
        →updating then stop the process and return w,b
        w,b = initialize_weights(x_train[0])
        N=len(x_train)
        tr_loss,ts_loss = [],[]
        #loss = OrderedDict()
        initial_loss = 0
        for epoch in tqdm(range(epochs)):
            #print(epoch)
            for i in range(x_train.shape[0]):
                dw = gradient_dw(x_train[i],y_train[i],w,b,alpha,N)
                db = gradient_db(x_train[i],y_train[i],w,b)
                w += (eta0*dw)
                b += (eta0*db)
            y_pred_tr,y_pred_ts = pred(w,b,x_train),pred(w,b,x_test)
            tr_loss.append(logloss(y_train,y_pred_tr))
            ts_loss.append(logloss(y_test,y_pred_ts))
            #loss[logloss(y_test,y_pred_ts)] = w
            #print(y_train,y_pred_tr)
            #print(y_test,y_pred_ts)
            #ts_loss = logloss(y_test,y_pred_ts)
            #if abs(initial_loss - ts_loss ) <= (10*(-3)):
            # print('(ts_loss - initial_loss):',(ts_loss, initial_loss),abs(ts_loss -
            →initial_loss),'\n')
            # break
            #print('(ts_loss - initial_loss):',(ts_loss, initial_loss),abs(ts_loss -
            →initial_loss),'\n')
            #initial_loss = ts_loss
        return w,b,tr_loss,ts_loss

```

[22]:

```

[23]: alpha=0.0001
eta0=0.0001
N=len(x_train)
epochs=14
w,b,tr_loss,ts_loss=train(x_train,y_train,x_test,y_test,epochs,alpha,eta0)

```

100%| 14/14 [00:20<00:00, 1.49s/it]

Goal of assignment

Compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of 10^{-3}

```

[24]: # these are the results we got after we implemented sgd and found the optimal
        →weights and intercept
w-clf.coef_, b-clf.intercept_

```




```
[24]: (array([[ -4.75139040e-03,  7.60245639e-03,  1.85102713e-03,  
             6.50362355e-05,  1.54498740e-03,  2.34086809e-03,  
            -9.09928936e-04,  2.16124544e-03,  5.21959720e-03,  
            -4.49834999e-03,  1.23628554e-03,  2.54417563e-03,  
             1.74962845e-03, -1.28756176e-03,  1.05365463e-03]]),  
      array([0.00279952]))
```

Plot epoch number vs train , test loss

- epoch number on X-axis
- loss on Y-axis

```
[25]: import matplotlib.pyplot as plt  
plt.title('Epoch number vs train , test loss')  
plt.plot(tr_loss,label='train-loss')  
plt.plot(ts_loss,label='test-loss')  
plt.legend()  
plt.show()
```



```
[25]:
```