

BackPropagation

There will be some functions that start with the word "grader" ex: grader_sigmod(), grader_forwardprop(), grader_backprop() etc, you should not change those function definition.

Every Grader function has to return True.

Loading data

```
In [ ]:

In [ ]:
import pickle
import numpy as np
from tqdm import tqdm
import matplotlib.pyplot as plt
path = '/content/drive/MyDrive/AI/Assignments/19.Backpropagation and Gradient Checking/practice/Copy of data.pkl'
with open(path, 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
print(X.shape, y.shape)

(595, 6)
(595, 5)

In [ ]:
```

Computational graph

- If you observe the graph, we are having input features [f1, f2, f3, f4, f5] and 9 weights [w1, w2, w3, w4, w5, w6, w7, w8, w9]
- The final output of this graph is a value L which is computed as $(Y-Y')^2$

Task 1: Implementing backpropagation and Gradient checking

Check this video for better understanding of the computational graphs and back propagation

```
In [ ]:
from IPython.display import YouTubeVideo
YouTubeVideo('1940vY66noo', width='1800', height='500')

Out [ ]:
```

- Write two functions
 - Forward propagation(Write your code in def forward_propagation())
- For easy debugging, we will break the computational graph into 3 parts.
- Part 1
- Part 2
- Part 3

```
def forward_propagation(X, y, W):
    # X: input data point, note that in this assignment you are having 5-d data points
    # y: output variable
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in graph, ..., W[8] corresponds to w9 in graph
    # you have to return the following variables
    # exp: part1 (compute the forward propagation until exp and then store the values in exp)
    # tanh: part2 (compute the forward propagation until tanh and then store the values in tanh)
    # sig = part3 (compute the forward propagation until sigmoid and then store the values in sig)
    # now compute remaining values from computational graph and get y'
    # write code to compute the value of L=(y-y')^2
    # compute derivative of L w.r.to y' and store it in dl
    # Create a dictionary to store all the intermediate values
    # store L, exp, tanh, sig, dl variables

    return (dictionary, which you might need to use for back propagation)

Backward propagation(Write your code in def backward_propagation())<b>

def backward_propagation(L, W, dictionary):
    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function
    # we need to compute the gradients of each weight [w1,w2,w3,...,w9]
    # Hint: you can use dict type to store the required variables
    # return dw, dw is a dictionary with gradients of all the weights

    return dw
```

Gradient clipping

Check this [blog](#) link for more details on Gradient clipping

we know that the derivative of any function is

$$\lim_{\epsilon \rightarrow 0} \frac{f(x+\epsilon) - f(x-\epsilon)}{2\epsilon}$$

- The definition above can be used as a numerical approximation of the derivative. Taking an epsilon small enough, the calculated approximation will have an error in the range of epsilon squared.
 - In other words, if epsilon is 0.001, the approximation will be off by 0.00001.
- Therefore, we can use this to approximate the gradient, and in turn make sure that backpropagation is implemented properly. This forms the basis of gradient checking!

Gradient checking example

lets understand the concept with a simple example: $f(w1, w2, x1, x2) = w1^2 \cdot x1 + w2 \cdot x2$

from the above function, lets assume $w1 = 1, w2 = 2, x1 = 3, x2 = 4$ the gradient of f.w.r.t $w1$ is

$$\frac{df}{dw1} = \frac{d}{dw1} (w1^2 \cdot x1 + w2 \cdot x2) = 2 \cdot w1 \cdot x1 = 2 \cdot 1 \cdot 3 = 6$$

let calculate the approximate gradient of $w1$ as mentioned in the above formula and considering $\epsilon = 0.0001$

$$\begin{aligned} d_{w1}^{approx} &= \frac{f(w1+\epsilon, x2, x1, x2) - f(w1-\epsilon, x2, x1, x2)}{2\epsilon} \\ &= \frac{(1+0.0001)^2 \cdot 3 + 2 \cdot 4 - ((1-0.0001)^2 \cdot 3 + 2 \cdot 4)}{2 \cdot 0.0001} \\ &= \frac{(1.00020001 \cdot 3 + 2 \cdot 4) - (0.99980001 \cdot 3 + 2 \cdot 4)}{0.0002} \\ &= \frac{(11.00060003 - 10.99940001)}{0.0002} \\ &= 5.999999999999 \end{aligned}$$

Then, we apply the following formula for gradient check: $gradient_check = \frac{|(dW) - (dW^{approx})|}{|(dW)| + |(dW^{approx})|}$

The equation above is basically the Euclidean distance normalized by the sum of the norm of the vectors. We use normalization in case that one of the vectors is very small. As a value for epsilon, we usually opt for $1e-7$. Therefore, if gradient check return a value less than $1e-7$, then it means that backpropagation was implemented correctly. Otherwise, there is potentially a mistake in your implementation. If the value exceeds $1e-3$, then you are sure that the code is not correct.

In our example: $gradient_check = \frac{(5.999999999999999 - 6)}{(5.999999999999999 + 6)} = 4.2514140356330737e-13$

you can mathematically derive the same thing like this

$$\begin{aligned} d_{w1}^{approx} &= \frac{f(w1+\epsilon, x2, x1, x2) - f(w1-\epsilon, x2, x1, x2)}{2\epsilon} \\ &= \frac{((w1+\epsilon)^2 \cdot x1 + w2 \cdot x2) - ((w1-\epsilon)^2 \cdot x1 + w2 \cdot x2)}{2\epsilon} \\ &= \frac{2 \cdot w1 \cdot x1}{2} \\ &= 2 \cdot w1 \cdot x1 \end{aligned}$$

Implement Gradient checking

(Write your code in def gradient_checking())

```
Algorithm

W = initialize randomly
def gradient_checking(data_point, W):

    # compute the L value using forward_propagation()
    # compute the gradients of W using backward_propagation()<font>
    approx_gradients = {}
    for each w1 weight value in W<font color='grey'>
        # add a small value to weight w1, and then find the values of L with the updated weights
        # subtract a small value to weight w1, and then find the values of L with the updated weights
        # compute the approximation gradients of weight w1<font color='grey'>
        approx_gradients.append(approximation gradients of weight w1)<font color='grey'>
    # compare the gradient of weights W from backward_propagation() with the aproximation gradients of weights with <b> gradient_check formula<font>
    return gradient_check<font>

NOTE: you can do sanity check by checking all the return values of gradient_checking(), they have to be zero. if not you have bug in your code
```

Task 2 : Optimizers

- As a part of this task, you will be implementing 3 type of optimizers(methods to update weight)
- Using the same computational graph that we mentioned above to do this task
- Initialize the 9 weights from normal distribution with mean=0 and std=0.01

Check below video and this blog

```
In [ ]:
from IPython.display import YouTubeVideo
YouTubeVideo('gtpoA1ayxa', width='1800', height='500')

Out [ ]:
```

```
Algorithm

for each epoch(1-100):
    for each data point in your data:
        using the Functions forward_propagation() and backward_propagation() compute the gradients of weights
        update the weights with help of gradients ex: w1 = w1-learning_rate*dw1

Implement below tasks<b>

Task 2.1: you will be implementing the above algorithm with Vanilla update of weights

Task 2.2: you will be implementing the above algorithm with Momentum update of weights

Task 2.3: you will be implementing the above algorithm with Adam update of weights

Note : If you get any assertion error while running grader functions, please print the variables in grader functions and check which variable is returning False .Recheck your logic for that variable .

In [141]:
from google.colab import drive
drive.mount('/content/drive')

Mounted at /content/drive

Forward propagation

In [142]:
def sigmoid(z):
    '''In this function, we will compute the sigmoid(z)'''
    # we can use this function in forward and backward propagation
    return 1/(1+np.exp(-z))

def forward_propagation(x, y, W):
    '''In this function, we will compute the forward propagation'''
    # x: input data point, note that in this assignment you are having 5-d data points
    # y: output variable
    # W: weight array, its of length 9, W[0] corresponds to w1 in graph, W[1] corresponds to w2 in graph, ..., W[8] corresponds to w9 in graph.
    # you have to return the following variables
    # exp: part1 (compute the forward propagation until exp and then store the values in exp)
    # tanh: part2 (compute the forward propagation until tanh and then store the values in tanh)
    # sig = part3 (compute the forward propagation until sigmoid and then store the values in sig)
    # now compute remaining values from computational graph and get y'
    # write code to compute the value of L=(y-y')^2
    # compute derivative of L w.r.to y' and store it in dl
    # Create a dictionary to store all the intermediate values
    # store L, exp, tanh, sig variables
    return (dictionary, which you might need to use for back propagation)

In [145]:
import pandas as pd
import math
import pickle
import numpy as np
from tqdm import tqdm
import matplotlib.pyplot as plt
path = '/content/drive/MyDrive/AI/Assignments/19.Backpropagation and Gradient Checking/practice/Copy of data.pkl'
with open(path, 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
print(X.shape, y.shape)

data = 9
dw = np.empty(shape=(data))
dx = pd.DataFrame(columns=['f1', 'f2', 'f3', 'f4', 'f5'])

In [146]:
def gradient(z):
    return 1/(1+np.exp(-z))

def grader_forwardprop(x, y, W):
    d = dict()
    p1 = np.exp([([W[0]*x[0]] + ([W[1]*x[1]]**2)*W[5])
    p2 = np.tanh([W[0]*x[0]] + ([W[1]*x[1]]**2)*W[5])
    d = np.gradient((np.sin(W[2]*x[2])) + ((W[3]*x[3]) + (W[4]*x[4])) + W[7] )
    y_hat = p2 + (p3*W[8])
    loss = (y - y_hat)**2
    dl = -2*(y - y_hat)
    d['exp'] = d1
    d['tanh'] = loss
    d['exp'] = p1
    d['tanh'] = p2
    d['sigmoid'] = p3
    return d

Grader function -1

In [147]:
def grader_sigmoid(z):
    val=sigmoid(z)
    assert(val==0.8887876779778823)
    return True
grader_sigmoid(2)

Out[147]:
True

Grader function -2

In [148]:
def grader_forwardprop(di):
    d1 = {data['dy_pr']:-1.9205278264819143}
    loss={data['loss']:-0.628804863872813}
    part1={data['exp']:-1.1272967040973583}
    part2={data['tanh']:-2.8479204329252446}
    part3={data['sigmoid']:-0.6279179387419721}
    assert(d1 and loss and part1 and part2 and part3)
    return True
w=np.ones(9)*0.1
di=forward_propagation(X[0],y[0],w)
grader_forwardprop(di)

Out[148]:
True

In [ ]:
```

Backward propagation

```
In [ ]:
def backward_propagation(L, W, dict):
    # L: the loss we calculated for the current point
    # dictionary: the outputs of the forward_propagation() function
    # write code to compute the gradients of each weight [w1,w2,w3,...,w9]
    # Hint: you can use dict type to store the required variables
    # dw = # in dw1 compute derivative of L w.r.to w1
    # dw2 = # in dw2 compute derivative of L w.r.to w2
    # dw3 = # in dw3 compute derivative of L w.r.to w3
    # dw4 = # in dw4 compute derivative of L w.r.to w4
    # dw5 = # in dw5 compute derivative of L w.r.to w5
    # dw6 = # in dw6 compute derivative of L w.r.to w6
    # dw7 = # in dw7 compute derivative of L w.r.to w7
    # dw8 = # in dw8 compute derivative of L w.r.to w8
    # dw9 = # in dw9 compute derivative of L w.r.to w9
    # return dw, dw is a dictionary with gradients of all the weights

def backward_propagation(x, w, d1):
    d = dict()
    dw['w1'] = float(d1['dy_pr']) * float(np.subtract(1, np.square(float(d1['tanh']))) * float((d1['exp'])** 2 * float(float(float(w[0]*x[0])+float(w[1]*x[1])**x[0]))
    dw['w2'] = float(d1['dy_pr']) * float(np.subtract(1, np.square(float(d1['tanh']))) * float((d1['exp'])** 2 * float(float(w[0]*x[0])+float(w[1]*x[1])**x[1]))
    dw['w3'] = float(d1['dy_pr']) * float(w[3]) * float(d1['sigmoid']) * float(np.subtract(1, float(d1['sigmoid']))) * float(x[3]) * float(np.sin(float(x[2]*w[2])))
    dw['w4'] = float(d1['dy_pr']) * float(w[4]) * float(d1['sigmoid']) * float(np.subtract(1, float(d1['sigmoid']))) * float(x[4]) * float(np.sin(float(x[2]*w[2])))
    dw['w5'] = float(d1['dy_pr']) * float(w[5]) * float(d1['sigmoid']) * float(np.subtract(1, float(d1['sigmoid']))) * float(d1['exp'])
    dw['w6'] = float(d1['dy_pr']) * float(np.subtract(1, np.square(float(d1['tanh']))) * float(d1['exp']))
    dw['w7'] = float(d1['dy_pr']) * float(w[7]) * float(d1['sigmoid']) * float(np.subtract(1, float(d1['sigmoid'])))
    dw['w8'] = float(d1['dy_pr']) * float(w[8]) * float(d1['sigmoid']) * float(np.subtract(1, float(d1['sigmoid'])))
    dw['w9'] = float(d1['dy_pr']) * float(d1['sigmoid'])

    return dw

Grader function -3

In [150]:
def grader_backprop(data):
    dw={data['dw1']:-0.2297332349870203}
    dw2={data['dw2']:-0.62140761671752055}
    dw3={data['dw3']:-0.09582540558626319}
    dw4={data['dw4']:-0.04485794122712423}
    dw5={data['dw5']:-0.0618077226486874549}
    dw6={data['dw6']:-0.6324751873437471}
    dw7={data['dw7']:-0.0618042650545231}
    dw8={data['dw8']:-0.04880228407315518}
    dw9={data['dw9']:-1.93184439687927}
    assert(dw1 and dw2 and dw3 and dw4 and dw5 and dw6 and dw7 and dw8 and dw9)
    return True
w=np.ones(9)*0.1
di=forward_propagation(X[0],y[0],w)
di=backward_propagation(X[0],w,d1)
grader_backprop(di)

Out[150]:
True

Implement gradient checking

In [ ]:

In [ ]:
W = initialize randomly
def gradient_checking(data_point, W):
    # compute the L value using forward_propagation()
    # compute the gradients of W using backward_propagation()
    approx_gradients = {}
    for each w1 weight value in W:
        # add a small value to weight w1, and then find the values of L with the updated weights
        # subtract a small value to weight w1, and then find the values of L with the updated weights
        # compute the approximation gradients of weight w1
        approx_gradients.append(approximation gradients of weight w1)
    # compare the gradient of weights W from backward_propagation() with the approximation gradients of weights with gradient_check formula
    return gradient_check

In [151]:
epsilon = float(0.0001)
from numpy.linalg import norm
def gradient_check(a,b):
    return (norm(a,b))/(norm(a)+norm(b))
w = np.random.randn(6)/np.sqrt(6)*0.1
di = forward_propagation(X[0],y[0],w)
dw = backward_propagation(X[0],y[0],w,d1)
approx_gradients = {}
for ind, val in enumerate(W):
    W[ind] = float(val) + epsilon
    f1 = forward_propagation(X[0],y[0],W)
    W[ind] = float(val) - epsilon
    f2 = forward_propagation(X[0],y[0],W)
    W[ind] = float(val)
    d1.approx = float((np.subtract(float(f1-'loss'), float(f2-'loss')))/(2*epsilon))
    approx_gradients.append(d1.approx)
for ind, val in enumerate(dw.keys()):
    print(ind, val)
    a = dw[val]
    b = approx_gradients[ind]
    res = gradient_check(a,b)
    print(val, res)

dw1 0.53317353905722e-07
dw2 0.91702086422196e-07
dw3 2.38427323901846e-11
dw4 3.39054662231628e-11
dw5 0.49523243971398e-11
dw6 2.474668958938e-07
dw7 3.217090663690535e-06
dw8 0.95044743096978e-10
dw9 1.31237859208499e-13

Task 2: Optimizers

for each epoch(1-100):
    for each data point in your data:
        using the Functions forward_propagation() and backward_propagation() compute the gradients of weights
        update the weights with help of gradients ex: w1 = w1-learning_rate*dw1

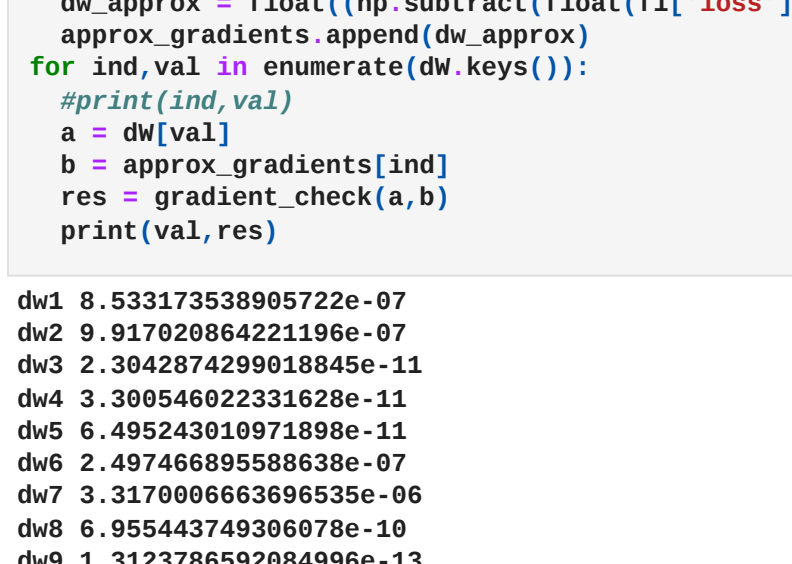
Algorithm with Vanilla update of weights

Task 2.1: you will be implementing the above algorithm with Vanilla update of weights

In [143]:
path = '/content/drive/MyDrive/AI/Assignments/19.Backpropagation and Gradient Checking/practice/Copy of data.pkl'
with open(path, 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
weights = np.random.normal(0,0.01,9)
lr = 0.001
plots = dict()
for epoch in range(1,101,1):
    for i in range(X.shape[0]):
        x = data[i, :5]
        y = data[i, -1]
        d1 = forward_propagation(x, y, weights)
        dw = np.array(list(backward_propagation(x, weights, d1).values()))
        weights = weights - lr * dw
        plots[epoch] = d1['loss']

https://stackoverflow.com/questions/27266341/plotting-a-python-dict-in-order-of-key-values
import matplotlib.pyplot as plt
lists = sorted(plots.items())
x, y = zip(*lists)
plt.plot(x, y, label='Vanilla Update')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.show()

In [ ]:
```




Algorithm with Momentum update of weights

- Task 2.2: you will be implementing the above algorithm with Momentum update of weights

```
In [142]:
path = '/content/drive/MyDrive/AI/Assignments/19.Backpropagation and Gradient Checking/practice/Copy of data.pkl'
with open(path, 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
guessa = 0
weights = np.random.normal(0,0.01,9)
lr = 0.001
plots = dict()
for epoch in range(1,101,1):
    for i in range(X.shape[0]):
        x = data[i, :5]
        y = data[i, -1]
        d1 = forward_propagation(x, y, weights)
        dw = np.array(list(backward_propagation(x, weights, d1).values()))
        momentum_term = ( guessa * w_previous
        gradient_term = ( lr * dw )
        new_weights = momentum_term + gradient_term
        guessa = w_previous + v_new
        plots[epoch] = d1['loss']

import matplotlib.pyplot as plt
lists = sorted(plots.items())
x, y = zip(*lists)
plt.plot(x, y, label='Momentum Update')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.show()

In [ ]:
```



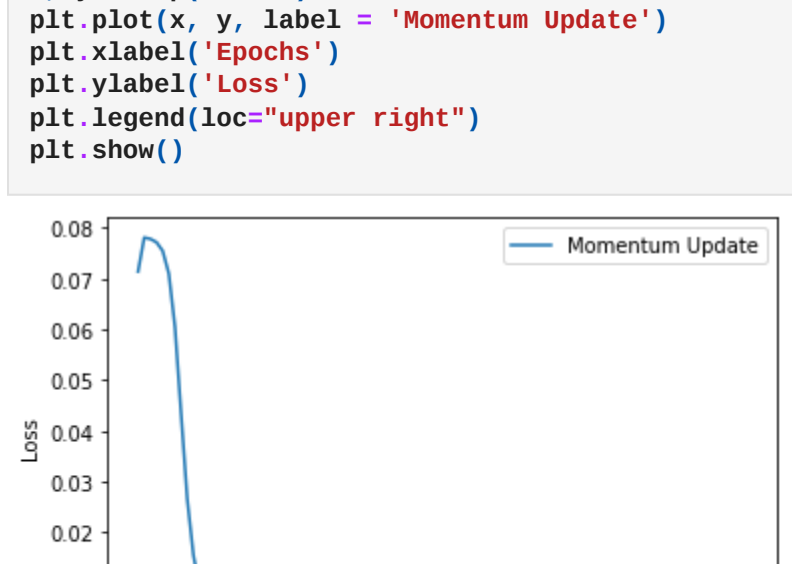
Algorithm with Adam update of weights

- Task 2.3: you will be implementing the above algorithm with Adam update of weights

```
In [141]:
path = '/content/drive/MyDrive/AI/Assignments/19.Backpropagation and Gradient Checking/practice/Copy of data.pkl'
with open(path, 'rb') as f:
    data = pickle.load(f)
print(data.shape)
X = data[:, :5]
y = data[:, -1]
betaz = 0.99
eps = 1e-8
weights = np.random.normal(0,0.01,9)
lr = 0.001
plots = dict()
for epoch in range(1,101,1):
    w = w_0
    for i in range(X.shape[0]):
        x = data[i, :5]
        y = data[i, -1]
        d1 = forward_propagation(x, y, weights)
        dw = np.array(list(backward_propagation(x, weights, d1).values()))
        m = ( beta1 * m ) + ( np.subtract(1.0 , betaz) ) * ( dw
        v = ( beta2 * v ) + ( np.subtract(1.0 , betaz) ) * ( np.square(dw)
        what = m / np.subtract(1.0, np.power(beta1,epoch))
        z = lr * ( what ) / (np.sqrt(what) + eps)
        weights = weights + z
        plots[epoch] = d1['loss']

import matplotlib.pyplot as plt
lists = sorted(plots.items())
x, y = zip(*lists)
plt.plot(x, y, label='Adam Update')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.show()

In [ ]:
```



Comparison plot between epochs and loss with different optimizers

```
In [144]:
import matplotlib.pyplot as plt
plt = sorted(plots.items())
p1 = sorted(plots.items())
p2 = sorted(plots.items())
p3 = sorted(plots.items())
a, b = zip(*p1)
c, d = zip(*p2)
e, f = zip(*p3)
plt.plot(c, d, label='Vanilla update')
plt.plot(e, f, label='Momentum update')
plt.plot(a, b, label='Adam update')
plt.xlabel('Epochs')
plt.ylabel('Loss')
plt.legend(loc='upper right')
plt.show()

In [ ]:
```

