

In this notebook, You will do amazon review classification with BERT.[Download data from [this link](#)]

It contains 5 parts as below. Detailed instructions are given in the each cell. please read every comment we have written.

1. Preprocessing
2. Creating a BERT model from the Tensorflow HUB.
3. Tokenization
4. getting the pretrained embedding Vector for a given review from the BERT.
5. Using the embedding data apply NN and classify the reviews.
6. Creating a Data pipeline for BERT Model.

instructions:

1. Don't change any Grader Functions. Don't manipulate any Grader functions. If you manipulate any, it will be considered as plagiarised.
2. Please read the instructions on the code cells and markdown cells. We will explain what to write.
3. please return outputs in the same format what we asked. Eg. Don't return List if we are asking for a numpy array.
4. Please read the external links that we are given so that you will learn the concept behind the code that you are writing.
5. We are giving instructions at each section if necessary, please follow them.

Every Grader function has to return True.

```
In [ ]: #in this assignment you need two files reviews.csv and tokenization file
#you can use gdown module to import both the files in colab from Google drive
#the syntax is for gdown is !gdown --id file_id
#please run the below cell to import the required files
```

```
In [ ]: # !gdown --id 1GsD8JlAc_0yJ-1151LNr6rLw83RRUPgt
# !gdown --id 13exfXiyiByluh1PfYK1EyZyizqxeCVG9
```

```
In [ ]: #pip install seaborn
```

```
In [42]: #all imports
import numpy as np
import pandas as pd
import tensorflow as tf
import tensorflow_hub as hub
from tensorflow.keras.models import Model
```

```
In [43]: tf.test.gpu_device_name()
```

```
Out[43]: '/device:GPU:0'
```

Grader function 1

```
In [44]: def grader_tf_version():
        assert((tf.__version__)>'2')
        return True
grader_tf_version()
```

```
Out[44]: True
```

Part-1: Preprocessing

```
In [45]: #Read the dataset - Amazon fine food reviews
reviews = pd.read_csv(r"Reviews.csv")
#check the info of the dataset
reviews.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 568454 entries, 0 to 568453
Data columns (total 10 columns):
#   Column                Non-Null Count  Dtype
---  -
0   Id                     568454 non-null  int64
1   ProductId              568454 non-null  object
2   UserId                  568454 non-null  object
3   ProfileName             568438 non-null  object
4   HelpfulnessNumerator    568454 non-null  int64
5   HelpfulnessDenominator  568454 non-null  int64
6   Score                   568454 non-null  int64
7   Time                     568454 non-null  int64
8   Summary                 568427 non-null  object
9   Text                     568454 non-null  object
dtypes: int64(5), object(5)
memory usage: 43.4+ MB

```

```

In [46]: #get only 2 columns - Text, Score
         #drop the NAN values

reviews = reviews[['Text','Score']]
reviews[reviews.isnull().any(axis=1)]
reviews.dropna(inplace=True)

```

```

In [47]: #if score> 3, set score = 1
         #if score<=2, set score = 0
         #if score == 3, remove the rows.

```

```

In [48]: ind = reviews[reviews['Score']==3].index
         reviews.drop(ind,axis=0,inplace=True)

reviews.loc[reviews['Score']<=2, "Score"] = 0
reviews.loc[reviews["Score"] >3, "Score"] = 1

```

```

In [49]: reviews.Score.value_counts()

```

```

Out[49]: 1    443777
         0    82037
         Name: Score, dtype: int64

```

Grader function 2

```

In [50]: reviews.shape

```

```

Out[50]: (525814, 2)

```

```

In [51]: def grader_reviews():
         temp_shape = (reviews.shape == (525814, 2)) and (reviews.Score.value_counts()[1]==443777)
         assert(temp_shape == True)
         return True
         grader_reviews()

```

```

Out[51]: True

```

```

In [52]: def get_wordlen(x):
         return len(x.split())
         reviews['len'] = reviews.Text.apply(get_wordlen)
         reviews = reviews[reviews.len<50]
         reviews = reviews.sample(n=100000, random_state=30)

```

```

In [53]: #remove HTML from the Text column and save in the Text column only

#fil = '<.*?>|&([a-z0-9]+|#[0-9]{1,6}|#x[0-9a-f]{1,6});'
#reviews['Text'] = reviews['Text'].str.replace(r'<[^>]*>', '', regex=True)
reviews['Text'] = reviews['Text'].str.replace(r'<.*?>|&([a-z0-9]+|#[0-9]{1,6}|#x[0-9a-f]{1,6});', '', regex=True)

```

```

In [54]: #print head 5
         reviews.head(5)

```

```

Out[54]:

```

	Text	Score	len
64117	The tea was of great quality and it tasted lik...	1	30
418112	My cat loves this. The pellets are nice and s...	1	31
357829	Great product. Does not completely get rid of ...	1	41
175872	This gum is my favorite! I would advise every...	1	27
178716	I also found out about this product because of...	1	22

```
In [55]: X = reviews[['Text', 'len']]
y = reviews['Score']
```


```
In [56]: #split the data into train and test data(20%) with Stratify sampling, random state 33,

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y,
                                                    stratify=y,
                                                    test_size=0.20,
                                                    random_state=33,
                                                    shuffle=True)
```


```
In [ ]:
```

```
In [16]: #plot bar graphs of y_train and y_test
```

```
import seaborn as sns
import matplotlib.pyplot as plt
y_train_ = pd.DataFrame(y_train.value_counts())
y_train_.plot.bar()
plt.xlabel("Count")
plt.ylabel("Y_train values")
plt.title("Distribution of Y_train")
plt.show()
```



```
In [17]: y_test_ = pd.DataFrame(y_test.value_counts())
y_test_.plot.bar()
plt.xlabel("Count")
plt.ylabel("Y_train values")
plt.title("Distribution of Y_test")
plt.show()
```



In [18]: reviews

Out[18]:

	Text	Score	len
64117	The tea was of great quality and it tasted lik...	1	30
418112	My cat loves this. The pellets are nice and s...	1	31
357829	Great product. Does not completely get rid of ...	1	41
175872	This gum is my favorite! I would advise every...	1	27
178716	I also found out about this product because of...	1	22
...
336657	Using this coffee and a stove top espresso mak...	1	39
498034	THE TASTE OF THIS M&M IS THE BEST. I USED IT I...	1	28
357766	Excellent Tea. I enjoy a cup every now and the...	1	21
326811	These oatmeal cookies have a great spice taste...	1	23
19261	This is the best coffee ever! I will never dri...	1	28

100000 rows × 3 columns

In [19]: `#saving to disk. if we need, we can load preprocessed data directly.`
`reviews.to_csv('preprocessed.csv', index=False)`

Part-2: Creating BERT Model

If you want to know more about BERT, You can watch live sessions on Transformers and BERT. we will strongly recommend you to read [Transformers](#), [BERT Paper](#) and, [This blog](#).

For this assignment, we are using [BERT uncased Base model](#). It uses L=12 hidden layers (i.e., Transformer blocks), a hidden size of H=768, and A=12 attention heads.

In [19]: `import os`
`## Loading the Pretrained Model from tensorflow HUB`
`tf.keras.backend.clear_session()`

`# maximum length of a seq in the data we have, for now i am making it as 55. You can change this`
`max_seq_length = 55`

`#BERT takes 3 inputs`

`#this is input words. Sequence of words represented as integers`
`input_word_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_word_ids")`

`#mask vector if you are padding anything`
`input_mask = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="input_mask")`

`#segment vectors. If you are giving only one sentence for the classification, total seg vector is 0.`

```

#If you are giving two sentences with [sep] token separated, first seq segment vectors are zeros and
#second seq segment vector are 1's
segment_ids = tf.keras.layers.Input(shape=(max_seq_length,), dtype=tf.int32, name="segment_ids")

#bert layer
#"https://tfhub.dev/tensorflow/bert\_en\_uncased\_L-12\_H-768\_A-12/4"
bert_layer = hub.KerasLayer("https://tfhub.dev/tensorflow/bert\_en\_uncased\_L-12\_H-768\_A-12/2", trainable=False)
#bert_layer = hub.KerasLayer("https://tfhub.dev/tensorflow/bert\_en\_uncased\_L-12\_H-768\_A-12/1", trainable=False)
pooled_output, sequence_output = bert_layer([input_word_ids, input_mask, segment_ids])

#Bert model
#We are using only pooled output not sequence out.
#If you want to know about those, please read https://www.kaggle.com/questions-and-answers/86510
bert_model = Model(inputs=[input_word_ids, input_mask, segment_ids], outputs=pooled_output)

```

In [20]: bert_model.summary()

Model: "model"

Layer (type)	Output Shape	Param #	Connected to
=====			
input_word_ids (InputLayer)	[(None, 55)]	0	

input_mask (InputLayer)	[(None, 55)]	0	

segment_ids (InputLayer)	[(None, 55)]	0	

keras_layer (KerasLayer)	[(None, 768), (None, 109482241		input_word_ids[0][0] input_mask[0][0] segment_ids[0][0]
=====			
Total params: 109,482,241			
Trainable params: 0			
Non-trainable params: 109,482,241			
=====			

In [21]: bert_model.output

Out[21]: <KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>

Part-3: Tokenization

In [22]: #getting Vocab file
vocab_file = bert_layer.resolved_object.vocab_file.asset_path.numpy()
do_lower_case = bert_layer.resolved_object.do_lower_case.numpy()

In [23]: import sys
sys.path.insert(0, "C:/Users/darsh/Downloads/Srujan/NLP_Transfer_learning_assignment/tokenization.py")
#pip install sentencepiece

In [24]: import tokenization #- #We have given tokenization.py file

In []:

In [26]: # Create tokenizer " Instantiate FullTokenizer"
name must be "tokenizer"
the FullTokenizer takes two parameters 1. vocab_file and 2. do_lower_case
we have created these in the above cell ex: FullTokenizer(vocab_file, do_lower_case)
please check the "tokenization.py" file the complete implementation

In [25]: # if you are getting error for sentencepiece module you can install it using below command while running this c
#!pip install sentencepiece
tokenizer=tokenization.FullTokenizer(vocab_file,do_lower_case)

Grader function 3

In [32]: #it has to give no error
def grader_tokenize(tokenizer):
 out = False
 try:
 out=('[CLS]' in tokenizer.vocab) and ('[SEP]' in tokenizer.vocab)
 except:
 out = False
 assert(out==True)
 return out
grader_tokenize(tokenizer)

Out[32]: True

```
In [33]: # Create train and test tokens (X_train_tokens, X_test_tokens) from (X_train, X_test) using Tokenizer and
# add '[CLS]' at start of the Tokens and '[SEP]' at the end of the tokens.

# maximum number of tokens is 55(We already given this to BERT layer above) so shape is (None, 55)

# if it is less than 55, add '[PAD]' token else truncate the tokens length.(similar to padding)

# Based on padding, create the mask for Train and Test ( 1 for real token, 0 for '[PAD]'),
# it will also same shape as input tokens (None, 55) save those in X_train_mask, X_test_mask

# Create a segment input for train and test. We are using only one sentence so all zeros. This shape will also
# type of all the above arrays should be numpy arrays

# after execution of this cell, you have to get
# X_train_tokens, X_train_mask, X_train_segment
# X_test_tokens, X_test_mask, X_test_segment
```

Example

[illegible]

```
In [34]: from tqdm import tqdm
X_train_tokens = list()
X_train_mask = list()
X_train_segment = list()
max_seq_length = 55
for i in tqdm(range(X_train.shape[0])):
    tokens = tokenizer.tokenize(X_train['Text'].values[i])
    tokens=tokens[0:(max_seq_length-2)]
    tokens=['[CLS]',*tokens,'[SEP]']
    if len(tokens)<max_seq_length:
        aa = ['[PAD]'] * (max_seq_length-len(tokens))
        tokens_=[*tokens] + aa
        X_train_mask.append([(1]*len(tokens)+[0]*(tokens_.count('[PAD]'))))
        tokens = tokenizer.convert_tokens_to_ids(tokens_)
```



```

if type(X_train_tokens) == np.ndarray:

    temp_shapes = (X_train_tokens.shape[1]==max_seq_length) and (X_train_mask.shape[1]==max_seq_length) and \
(X_train_segment.shape[1]==max_seq_length)
    #print('temp_shapes',temp_shapes)

    segment_temp = not np.any(X_train_segment)
    #print('segment_temp',segment_temp)

    mask_temp = np.sum(X_train_mask==0) == np.sum(X_train_tokens==0)
    #print('mask_temp',mask_temp)

    no_cls = np.sum(X_train_tokens==tokenizer.vocab['[CLS]'])==X_train_tokens.shape[0]
    #print('no_cls',no_cls)

    no_sep = np.sum(X_train_tokens==tokenizer.vocab['[SEP]'])==X_train_tokens.shape[0]
    #print('no_sep',no_sep)

    out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep
    #print('out',out)

else:
    print('Type of all above token arrays should be numpy array not list')
    out = False
assert(out==True)
return out

grader_alltokens_train()

```

Out[40]: True

Grader function 5

```

In [41]: def grader_alltokens_test():
out = False
if type(X_test_tokens) == np.ndarray:

    temp_shapes = (X_test_tokens.shape[1]==max_seq_length) and (X_test_mask.shape[1]==max_seq_length) and \
(X_test_segment.shape[1]==max_seq_length)

    segment_temp = not np.any(X_test_segment)

    mask_temp = np.sum(X_test_mask==0) == np.sum(X_test_tokens==0)

    no_cls = np.sum(X_test_tokens==tokenizer.vocab['[CLS]'])==X_test_tokens.shape[0]

    no_sep = np.sum(X_test_tokens==tokenizer.vocab['[SEP]'])==X_test_tokens.shape[0]

    out = temp_shapes and segment_temp and mask_temp and no_cls and no_sep

else:
    print('Type of all above token arrays should be numpy array not list')
    out = False
assert(out==True)
return out
grader_alltokens_test()

```

Out[41]: True

Part-4: Getting Embeddings from BERT Model

We already created the BERT model in the part-2 and input data in the part-3. We will utilize those two and will get the embeddings for each sentence in the Train and test data.

In [42]: bert_model.input

Out[42]: [<KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'input_word_ids')>, <KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'input_mask')>, <KerasTensor: shape=(None, 55) dtype=int32 (created by layer 'segment_ids')>]

In [43]: bert_model.output

Out[43]: <KerasTensor: shape=(None, 768) dtype=float32 (created by layer 'keras_layer')>


```
In [44]: # get the train output, BERT model will give one output so save in
# X_train_pooled_output
#this cell will take some time to execute, make sure that you have stable internet connection
X_train_pooled_output=bert_model.predict([X_train_tokens,X_train_mask,X_train_segment])
```

```
In [45]: # get the test output, BERT model will give one output so save in
# X_test_pooled_output
X_test_pooled_output=bert_model.predict([X_test_tokens,X_test_mask,X_test_segment])
```

```
In [46]: ##save all your results to disk so that, no need to run all again.
pickle.dump((X_train_pooled_output, X_test_pooled_output),open('final_output.pkl','wb'))
```

```
In [26]: import pickle
```

```
In [27]: X_train_pooled_output, X_test_pooled_output= pickle.load(open('final_output.pkl', 'rb'))
```

```
In [ ]:
```

Grader function 6

```
In [49]: #now we have X_train_pooled_output, y_train
#X_test_pooled_output, y_test

#please use this grader to evaluate
def grader_output():
    assert(X_train_pooled_output.shape[1]==768)
    assert(len(y_train)==len(X_train_pooled_output))
    assert(X_test_pooled_output.shape[1]==768)
    assert(len(y_test)==len(X_test_pooled_output))
    assert(len(y_train.shape)==1)
    assert(len(X_train_pooled_output.shape)==2)
    assert(len(y_test.shape)==1)
    assert(len(X_test_pooled_output.shape)==2)
    return True
grader_output()
```

```
Out[49]: True
```

Part-5: Training a NN with 768 features

Create a NN and train the NN.

1. You have to use AUC as metric. Do not use `tf.keras.metrics.AUC`
You have to write custom code for AUC and print it at the end of each epoch
2. You can use any architecture you want.
3. You have to use tensorboard to log all your metrics and Losses. You have to send those logs.
4. Print the loss and metric at every epoch.
5. You have to submit without overfitting and underfitting.

```
In [35]: ##imports
from tensorflow.keras.layers import Input, Dense, Activation, Dropout, LSTM, Conv1D, Flatten, Conv2D, Embedding
from tensorflow.keras.models import Model, Sequential
```

```
In [36]: ##create an Neural Network and train your model on X_train_pooled_output and y_train
# you can start as follows
input_layer=Input(shape=(X_train_pooled_output.shape[1],))
```

```
In [37]: import os
import datetime
import tensorflow as tf
log_dir = os.path.join("logs", 'fits', datetime.datetime.now().strftime("%Y%m%d-%H%M%S"))
tensorboard_callback = tf.keras.callbacks.TensorBoard(log_dir=log_dir,histogram_freq=1,write_graph=True)
```

```
In [38]: %load_ext tensorboard
```

```
In [39]: checkpoint_path = "C:/Users/darsh/Downloads/Srujan/NLP_Transfer_learning_assignment/cpp1-{epoch:04d}.ckpt"
checkpoint_dir = os.path.dirname(checkpoint_path)
```

```
In [57]: xtrain=X_train_pooled_output.reshape(-1,768,1)
xtest=X_test_pooled_output.reshape(-1,768,1)

ytrain = tf.keras.utils.to_categorical(y_train.values, num_classes=2)
ytest= tf.keras.utils.to_categorical(y_test.values, num_classes=2)
```

```
In [ ]:
```

```
In [58]: model1 = Sequential()
input_layer = Input(shape=(xtrain.shape[1],xtest.shape[2]))
model1.add(input_layer)
model1.add(Conv1D(128, 3, activation='relu',padding='same')) # input_shape = (768,1)
model1.add(Conv1D(256, 3, activation='relu',padding='same'))
model1.add(Conv1D(256, 3, activation='relu',padding='same'))
# flat
model1.add(Flatten())

model1.add(Dense(2, activation='softmax'))
model1.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv1d (Conv1D)	(None, 768, 128)	512
conv1d_1 (Conv1D)	(None, 768, 256)	98560
conv1d_2 (Conv1D)	(None, 768, 256)	196864
flatten (Flatten)	(None, 196608)	0
dense (Dense)	(None, 2)	393218
Total params: 689,154		
Trainable params: 689,154		
Non-trainable params: 0		

```
In [59]: import tensorflow as tf
from sklearn.metrics import roc_auc_score

def auc(y_true, y_pred):
    return tf.numpy_function(roc_auc_score, (y_true, y_pred), tf.double)
```

In []:

In []:

```
In [60]: from tensorflow.keras.optimizers import SGD,Adam
opt = Adam()

callbacks = [
    tf.keras.callbacks.ModelCheckpoint('./NLP_TL_Best_model.h5', save_weights_only=False, save_best_only=True,
                                       mode='min', monitor='loss',verbose=1,save_freq='epoch'),
    tf.keras.callbacks.ReduceLROnPlateau(monitor='loss', patience=1,mode='min',verbose=1),
    tensorboard_callback
]
model1.compile(
    loss='categorical_crossentropy',
    optimizer="adam",
    metrics=["accuracy",auc]
)
```

In []:

```
In [61]: model1.fit(xtrain,ytrain,batch_size=128,
                    validation_data=(xtest,ytest),callbacks=callbacks,epochs=5,verbose=1)
```

```
Epoch 1/5
625/625 [=====] - 66s 88ms/step - loss: 0.2486 - accuracy: 0.9007 - auc: 0.8968 - val_
loss: 0.1938 - val_accuracy: 0.9201 - val_auc: 0.9412

Epoch 00001: loss improved from inf to 0.24860, saving model to .\NLP_TL_Best_model.h5
Epoch 2/5
625/625 [=====] - 70s 113ms/step - loss: 0.1925 - accuracy: 0.9222 - auc: 0.9490 - val
_loss: 0.1861 - val_accuracy: 0.9283 - val_auc: 0.9459

Epoch 00002: loss improved from 0.24860 to 0.19253, saving model to .\NLP_TL_Best_model.h5
Epoch 3/5
625/625 [=====] - 67s 108ms/step - loss: 0.1866 - accuracy: 0.9242 - auc: 0.9529 - val
_loss: 0.1868 - val_accuracy: 0.9262 - val_auc: 0.9487

Epoch 00003: loss improved from 0.19253 to 0.18660, saving model to .\NLP_TL_Best_model.h5
Epoch 4/5
625/625 [=====] - 67s 107ms/step - loss: 0.1790 - accuracy: 0.9274 - auc: 0.9553 - val
_loss: 0.1884 - val_accuracy: 0.9255 - val_auc: 0.9467

Epoch 00004: loss improved from 0.18660 to 0.17902, saving model to .\NLP_TL_Best_model.h5
Epoch 5/5
625/625 [=====] - 67s 107ms/step - loss: 0.1781 - accuracy: 0.9286 - auc: 0.9570 - val
_loss: 0.1760 - val_accuracy: 0.9318 - val_auc: 0.9506

Epoch 00005: loss improved from 0.17902 to 0.17806, saving model to .\NLP_TL_Best_model.h5
```

```
Out[61]: <keras.callbacks.History at 0x225bedafee0>
```

```
In [ ]:
```

```
In [62]: %tensorboard --logdir "C:/Users/darsh/Downloads/Srujan/NLP_Transfer_learning_assignment/logs/fits/20221201-04221"
```

Connection refused

Failed to load URL <https://html2pdf.com:6006/>.

QtNetwork Error 1

```
In [ ]:
```

```
In [ ]:
```

1. Pipeline is a way to codify and automate the workflow.
2. Download the test.csv file from here [here](#)

```
#there is an alterante way to load files from Google drive directly to your Colab session
# you can use gdown module to import the files as follows
#for example for test.csv you can write your code as !gdown --id file_id (remove the # from next line and run i
```

```
#read the csv file
test_df= pd.read_csv('test.csv')
```

1. You have to write a function that takes the test_df, trained model and the required parameters as input.
 2. Perform all the preprocessing steps inside the function.
 - Remove all the html tags
 - Now do tokenization [Part 3 as mentioned above]
 - Create tokens, mask array and segment array
 - Get Embeddings from BERT Model [Part 4 as mentioned above] , let it be X_test
 - Print the shape of output(X_test.shape). You should get (352, 768)
 3. Predict the output of X_test with the neural network model which we trained earlier.
 4. Return the occurrences of class labels from the function.
- The output should be the count of datapoints classified as 1 or 0.

```
test_df['Text'] = test_df['Text'].str.replace(r'<.*?>|&([a-z0-9]+|#[0-9]{1,6}|#x[0-9a-f]{1,6});', '', regex=True)
```

```

from tqdm import tqdm
Xtest_tokens = list()
Xtest_mask = list()
Xtest_segment = list()
#max_seq_length = 55
for i in tqdm(range(test_df.shape[0])):
    tokens = tokenizer.tokenize(test_df['Text'].values[i])
    tokens=tokens[0:(max_seq_length-2)]
    tokens=['[CLS]','*tokens','[SEP]']
    if len(tokens)<max_seq_length:
        aa = ['[PAD]'] * (max_seq_length-len(tokens))
        tokens =[*tokens] + aa
        Xtest_mask.append(([1]*len(tokens)+[0]*(tokens_.count('[PAD]'))))
        tokens = tokenizer.convert_tokens_to_ids(tokens_)
        #print("<55",len(tokens_),tokens_.count('[PAD]'),tokens_,'\n')
    else:
        #print('>=55',len(tokens))
        Xtest_mask.append([1]*len(tokens))
        tokens = tokenizer.convert_tokens_to_ids(tokens)

    Xtest_tokens.append(tokens)
    seg = [0 for i in range(len(tokens))]
    Xtest_segment.append(seg)

#print(np.array(Xtest_tokens).shape,np.array(Xtest_mask).shape,np.array(Xtest_segment).shape)
Xtest_tokens = np.array(Xtest_tokens)
Xtest_mask = np.array(Xtest_mask)
Xtest_segment = np.array(Xtest_segment)
X_test_=bert_model.predict([Xtest_tokens,Xtest_mask,Xtest_segment])

```

```
X_test_.shape
```

(352, 768)

```
pred = (model1.predict(X_test_.reshape(-1,768,1)))
```

```
out = list()
for i in range(pred.shape[0]):
    out.append(np.argmax(pred[i]))
```

```
Out[77]: (29, 323)
```

In []:

```
100% |██████████| 352/352 [00:00<00:00, 1691  
.66it/s]
```

```
Out[68]: (39, 313)
```

Please write your observations at the end of notebook and explain each and every step you followed in solving this assignment.

1. Preprocessing:

- Downloaded the amazon reviews data and only Text, Score features were considered.
- Reviews which have a score greater than 3 is considered as a positive class or `class_label=1`
- Reviews that have a score less than or equal to 2 are considered as a negative class or `class_label=0`
- Reviews that have a score equal to 3 are dropped
- From this new dataset with features = [Text, Score, Class_label], a random sample of 100000 reviews are selected whose word length is less than 50.

2. Creating a BERT model from the Tensorflow HUB:

- We have created this model to get word embeddings that are used in the model.
- BERT model takes 3-inputs i.e [input word ids, Segment ids,input mask]

3. Tokenization:

- a. To generate these inputs, we use the tokenization.py file and generate these 3-inputs
- b. Preprocessing of Text like creating input ids,input mask,segment ids and padding also

applied to obtain

Text of size equals `max_sequence_length`. This is done because we can send a batch of rows at a time.

4. Getting the pre-trained embedding Vector for a given review from the BERT.

a. Now all these 3 inputs are sent to the BERT model to get the word embeddings of size 768 dimensions for all rows in data.

5. Using the embedding data apply NN and classify the reviews:

a. Using these word embeddings from the BERT model, a neural network with 3-Conv1D layers and a Dense layer is created.

b. The input layer for this neural network is modified

c. Train and test datasets were created and `class_label` is converted to categorical type.

d. Adam optimizer, custom AUC metric, `categorical_crossentropy` as a loss were considered for the neural network model

e. Model has been trained for 5 epochs with `batch_size=128`.

6. Creating a Data Pipeline for BERT Model:

a. created a data pipeline that does all the above steps and returns the counts of each label.

In []:

Loading [MathJax]/jax/output/CommonHTML/fonts/TeX/fontdata.js